

Incident Report — INC-2024-047 (PRELIMINARY — POST-MORTEM NOT WRITTEN)

Status: OPEN — Root cause investigation incomplete

Incident ID: INC-2024-047

Severity: P1 — Production down

Date: 14 May 2026

Start time: 02:17 UTC

End time: 03:04 UTC

Duration: 47 minutes

Detected by: Client call (Meridian Finance operations team)

Resolved by: Jordan Reeves (on-call), manual pod restart

Impact

Metric | Value

Services affected | booking-api, payment-api

Clients affected | All 140+ clients

Booking requests failed | ~3,200 (estimated)

Payment attempts failed | ~890 (estimated)

SLA breach | YES — 47 min exceeds 30-min SLA

SLA penalty issued | £18,000

Churn risk triggered | Meridian Finance, Crestfield Retail

Timeline

Time (UTC) | Event

02:17 | First error logs in booking-api (not alerted — no monitoring)

02:19 | payment-api begins returning 500s

02:23 | SQS queue depth begins spiking — notification backlog grows

02:31 | auth-service readiness probe begins failing

~03:00 | Meridian Finance operations team notices booking failures

03:04 | Meridian Finance calls NimbusCloud emergency line

03:09 | Jordan Reeves paged — manual investigation begins

03:22 | Root cause identified as suspected: Kubernetes pod memory exhaustion

03:24 | `kubectl rollout restart deployment/booking-api`

03:26 | `kubectl rollout restart deployment/payment-api`

03:28 | Pods back to Running state

03:30 | Services recovering — error rate dropping

03:47 | All services confirmed healthy

03:52 | Meridian Finance updated — incident closed

Preliminary Root Cause

Suspected: Kubernetes pod memory limit hit — pods OOMKilled, entered CrashLoopBackOff

Actual root cause: NOT CONFIRMED — no logs captured during outage window

Evidence gap: Prometheus had no scrape targets — no metrics from 02:00–03:50 UTC

■ ■ The actual root cause remains UNKNOWN because monitoring was not in place.

This post-mortem is INCOMPLETE. A full post-mortem must be written.

Assigned to: NEW CLOUD ENGINEER (you) — Phase 6

Outstanding Actions

| Action | Owner | Due | Status

- 1 | Write full post-mortem with confirmed root cause | Incoming cloud engineer | Phase 6 | OPEN
 - 2 | Implement Prometheus scrape targets | Incoming cloud engineer | Phase 5 | OPEN
 - 3 | Create Grafana dashboards for all 4 services | Incoming cloud engineer | Phase 5 | OPEN
 - 4 | Configure CloudWatch alarms for latency and errors | Incoming cloud engineer | Phase 5 | OPEN
 - 5 | Fix auth-service readiness probe | Incoming cloud engineer | Phase 2 | OPEN
 - 6 | Document memory limit recommendations | Incoming cloud engineer | Phase 7 | OPEN
-

Notes from Jordan Reeves (outgoing engineer)

"I didn't have time to set up proper monitoring before the incident. The pod likely hit its memory limit — I've seen it happen before on high-traffic nights. The fix was a rollout restart but I can't prove it. The real problem is we have no visibility into what happened. Whoever inherits this needs to sort out

Prometheus and Grafana before the next incident — not after it."

— Jordan Reeves, 2026-05-01